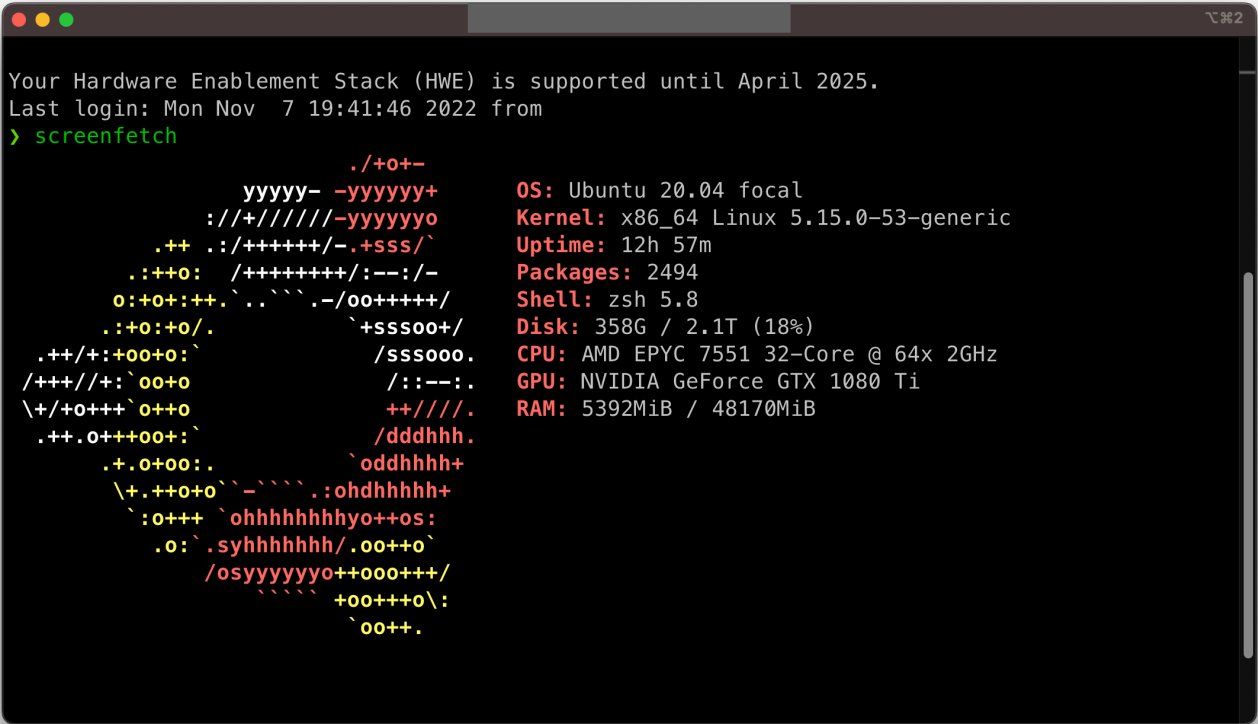


CSC 3150 Assignment 4 Report

Tianhao Shi, 120090472

Nov, 2022

Environment



CUDA Version: 11.6

```
> nvidia-smi
```

Wed Nov 9 20:18:00 2022

NVIDIA-SMI 510.85.02		Driver Version: 510.85.02				CUDA Version: 11.6			

GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC		
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.		
=====									
0	NVIDIA GeForce ...	On	00000000:41:00.0	Off			N/A		
23%	27C	P8	8W / 250W	52MiB / 11264MiB	0%	Default	N/A		

Processes:									
GPU	GI	CI	PID	Type	Process name	GPU Memory			
	ID	ID				Usage			
=====									
0	N/A	N/A	2112	G	/usr/lib/xorg/Xorg	49MiB			

Also tested on the provided cluster

Execution

Method 1: Use makefile

```
cd Assignment_3_120090472/source
make #(compile+run)
# make build #(build only)
```

Method 2: Batch script

```
cd Assignment_3_120090472/source
bash ./slurm.sh
```

Overview of Project

This project simulates a basic file system using the GPU memory. There is only one volume, which is further divided into three parts:

1. 4KB Volume Control Block
2. 32KB File Control Blocks (32B*1024)
3. 1MB Space for storing the actual component of the file content.

Aside from the main volume, there is also a 128B temporary space for global variables. This file system is created and initialized by launching a single thread ($\lll 1, 1 \ggg$) with an input buffer and a output buffer in the main file (`main.cu`). The input buffer stores all the data to be written to the file system, while all the read operations stores data in the output buffer. The set of operation on the file system is defined in the `user_program()` function within the `user_program.cu` file. The whole file system volume can be viewed as a very large 1-d uchar array, and all the operations are manipulating the array.

Design of key components

Volume Control Block (VCB)

In this file system, each file occupies some bytes. However, to manage files easily, we group every 32 bytes of the file as a block. When allocating space for files, we always allocate entire blocks to files, even if the file size is not a multiple of 32 (which means this file system has **internal fragmentation**). To keep track of the free space on the disk, we use the volume control block. This block serves like a 'snapshot' of the actual storage area.

The underlying data structure of the VCB is a bitmap. In my implementation, each byte of the VCB represents the status of 8 blocks. 1 means the block is occupied, while 0 means the block is available for allocation. Since the VCB area is 4KB, there are $4K \times 8 = 32K$ blocks. These blocks can store $32KB \times 32KB = 1MB$ data, which corresponds to the size of the storage area for data.

File Control Blocks (FCB)

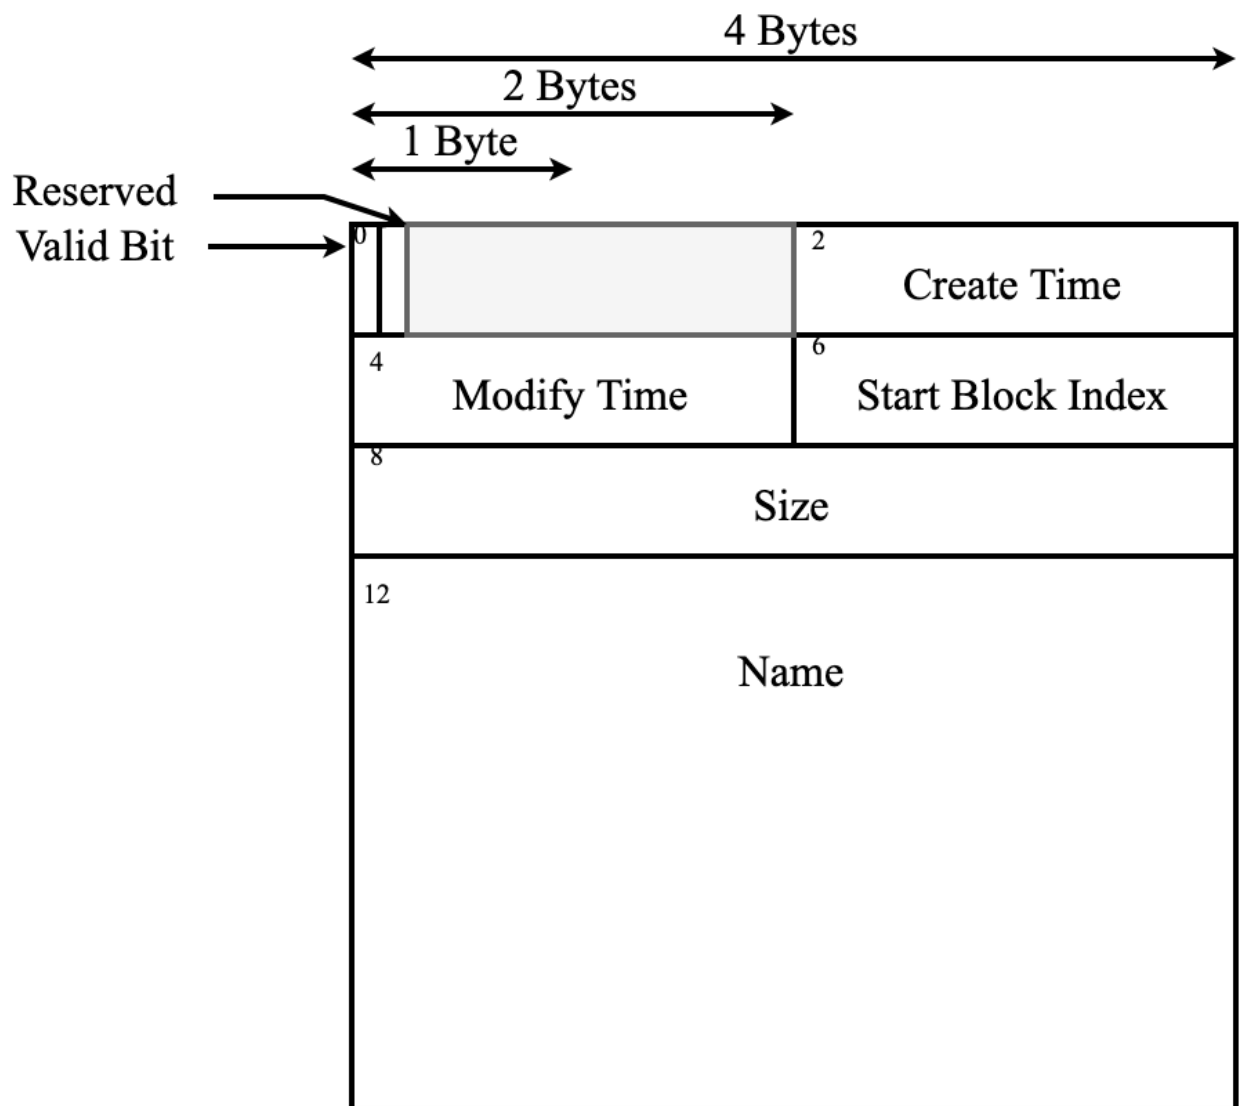
As mentioned above, the VCB only records the availability of the disk space. To identify different files, we also need a place to store auxiliary information about files. The information about a file is called its **metadata**. We use 32 Bytes to record various information of a file. Each of these 32 Byte space is called a **File Control Block**. The size of FCB is consistent for every file, and the type of information stored in a particular area within the FCB is also the same for different files.

Source

In the basic part of the project, each FCB stores the following information:

- Valid Bit: 1 bit, indicates whether this FCB is valid. If it's cleared, the rest attributes will not be valid.
- Create Time: 2 bytes, indicates the creation time of a file. Set when the file is created and never changed.
- Modify Time: 2 bytes, indicates the most recent update time of the file. Set when: a) the file is created; b) when the file is updated.
- Start Block Index: 2 bytes, indicates the index of the starting block of the file. Using this information, we can acquire the base address of a file. Set whenever the file is write. May change whenever the file is allocated to a different location (to be discussed later in `fs_write()`).
- Size: 4 bytes, indicates the actual size of the file. When the file is created, it is set to 0. Later its value will be set whenever the file is written. I use 4 bytes because the maximum size of the file is 1MB (this is the case when there is only 1 file occupying the whole storage area). Although using 3 bytes is also enough, since $3 \times 8 > 20$, to guarantee address alignness, I still use 4 bytes.
- Name: 20 bytes (may be smaller), indicates the actual name of the file. Note that whenever we set or read the file, we use a `char*` pointer, which retrives until the first `/0`.

The following graph illustrates how I use the 32 bytes of the FCB. Note that for the sake of illustration, I showed FCB as a 2D array, but in fact it should be an 1d array.



Bonus

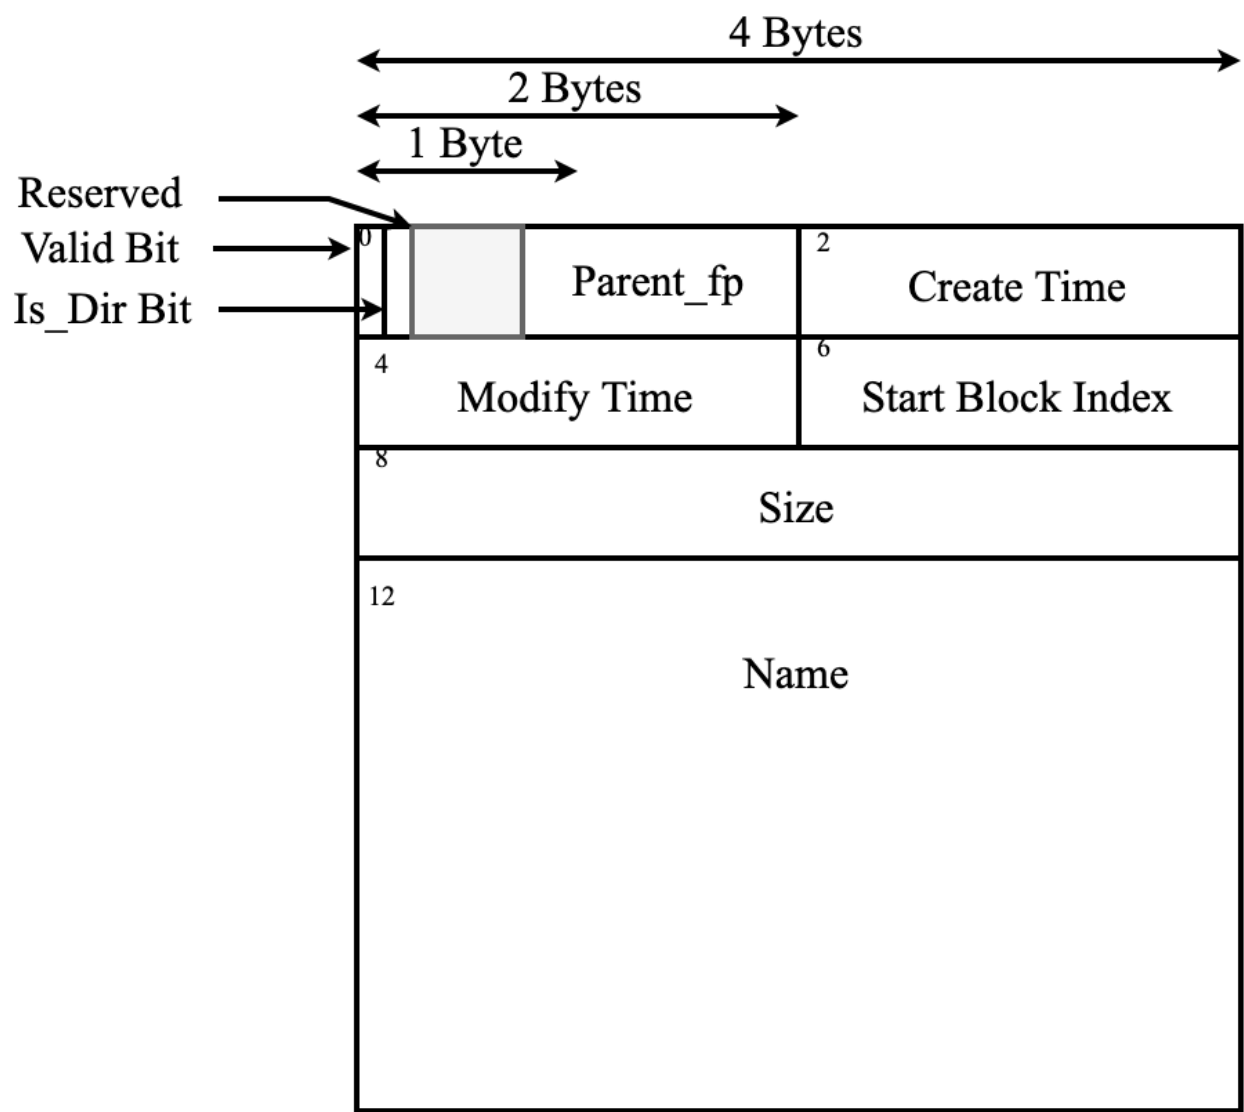
It should be noticed that in the basic version of the project there is no subdirectory, and every file is stored under the root directory. However, in the bonus part we have to design the tree-structure. This indicates there is always 1 FCB used to record the metadata of the root directory (in my case, I use FCB[0]). Moreover, introducing subdirectories also poses the following challenge:

1. How to differentiate a file from a folder(directory)?
2. How to identify the parent of a file/directory?

To tackle these challenges, I introduced two attributes to the FCB:

1. `is_dir` bit (1 bit): indicates whether this is a file or a folder. 1 means this is a folder, and 0 means this is a file. This bit limits operation allowed on this FCB, as I will discuss later.
2. `parent_fp` (10 bits): indicates the position of a folder/directory's parent directory. Basically it's the index of that directory's FCB

In the actual operation, these two attributes, together with the valid bit is always retrived together by taking 2 bytes (MISC_ATTR_LENGTH) starting from 0th byte. The actual value of each attribute is acquired through bit operations. The updated FCB is illustraed as follows:



File Pointers

The file pointer contains information about the FCB index of a file. In the user level, i.e, when passed between provided API functions like fs_open()/fs_read()/fs_write(), its least significant bit in binary representation states the type of that file pointer. There are two types of file pointers in the user leve:

1. Read file pointer: LSB is 0
2. Write file pointer: LSB is 1

When passing a file pointer to an incompatable function (e.g, passing a write file pointer to fs_read()) will cause an error.

Aside from the LSB, the other bits are simply the index of the file's FCB among the 1024 FCB entries. This means its value, if valid ranges between 0 and 1023. When it equals to 1024, it indicates the requested file is not present in the FCB table.

When the fp is not passed in the user level, its type is not important anymore, so we omit the type bit. I.e., when not passed to fs_read()/fs_write(), the file pointer value equals to the index of the file's FCB entry within the FCB table.

Content of the files

The content of the file is stored starting from the base address (4K+32K=36K). Each byte of the file corresponds to a ucar within the volume. The base address, as described earlier, can be acquired using the start byte:

$$start_address = storage_base_addr + block_idx \times block_size$$

and the size of the file can also be acquired from the FCB's metadata.

Temporary Usage

Source

In the basic version of this project, 12/128 bytes of temporary usage space is utilized to store 3 global variables:

- gtime: the current time of the file system. Note that I assume there will be less than 2^{16} operations in the file system. Also, in my design, each command contributes to at most 1 increment in the current time. The details will be discussed in the function design part. In this system, the time is rigorously increasing. On initialization it is set to 0.
- gfilenum: the current number of file within the file system. On initialization, it is set to 0.
- glastblock: the index of next possibly available block. This variable is used because I use **next fit algorithm** to allocate space for new file (or to move old file to a new position). On initialization, it is set to 0, because the 0th block is available.

Bonus

In the bonus part, 20/128 bytes of the temporary space is used. Aside from the 3 global variables, I add two more variables:

- gcwd: tracks the current directory's FCB index. On initialization, it is set to 0, indicating a filesystem is in the root directory.
- glevel: tracks the depth of current directory. the root directory is located in the 0th level. Every file within the root directory is in the 1st level. Therefore, on initialization, this variable is set to 1.

Design of Functions

`fs_init()`

This function is used to initialize the file system. Aside from the code provided in the template, I add the following codes:

source

1. Clear all the bytes in the VCB area.
2. Invalidate all the FCB entries.

bonus

1. Create the root directory's FCB using the `fs_gsys(MKDIR)`
2. It should be noted that the time after `fs_init` should be 0, but creating the root FCB increases the system time. Therefore, we should reset the system time.

`fs_open()`

This function receives a file name and returns a user-level file pointer(LSB indicates mode).

source

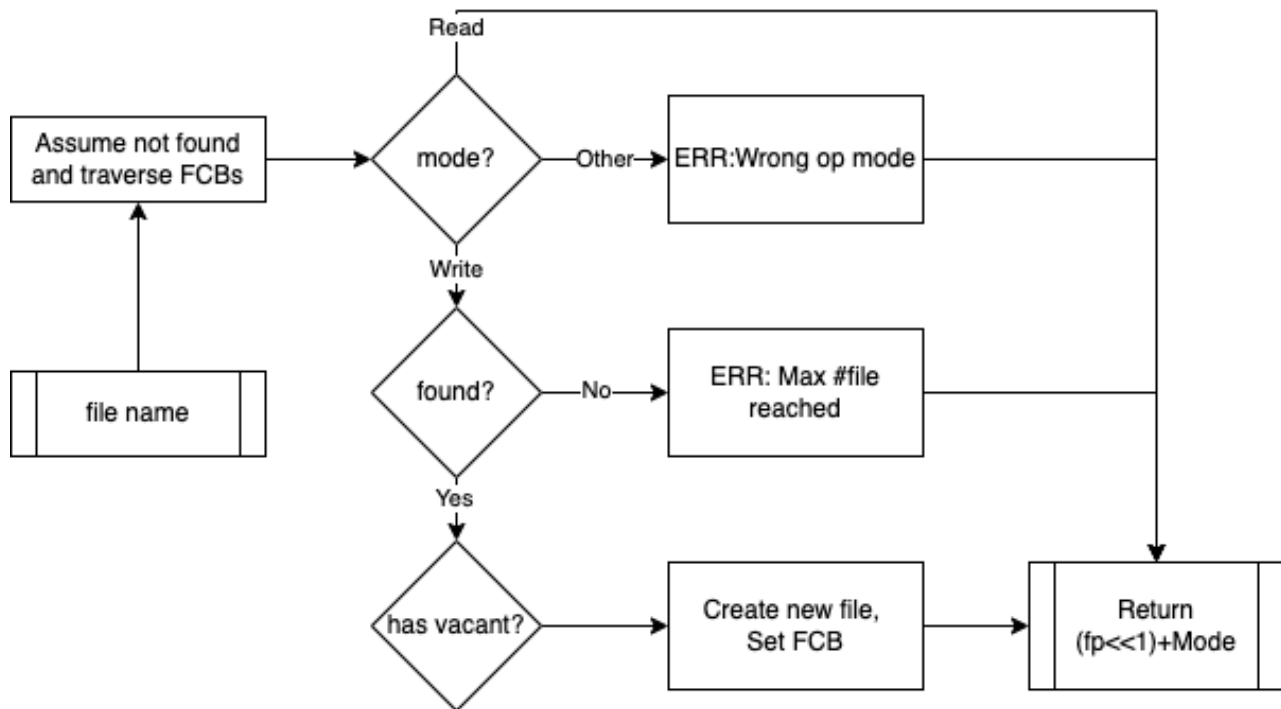
First, the function will traverse the FCB table until:

1. It found a FCB with name same as the requested name
2. It traversed all valid FCB and found a vacant FCB

When the requested filename is not present:

1. If the file is opened in *read* mode, an error is raised
2. If the file is opened in *write* mode
 1. If there is a vacant FCB, it will create a new file. This increases the file number and the system time
 - The *valid bit* is set
 - The *size* is set as 0
 - The *create time* and *modify time* is set as the current system time (before increment)
 - The *name* is set as the requested
 - No block is allocated on creation
 2. Else, it means maximum number of file is reached, and an error is raised

On the other hand, if the filename is present, this function simply left shifts the found FCB index by 1 bit and append the mode of the file pointer.



bonus

The bonus implementation of `fs_open()` is almost identical with the basic version. Except that when searching for the filenames in the FCBs, we have to make sure the filename and the the parent_fp are both the same. Since a file with same name could appear in different directory.

fs_write()

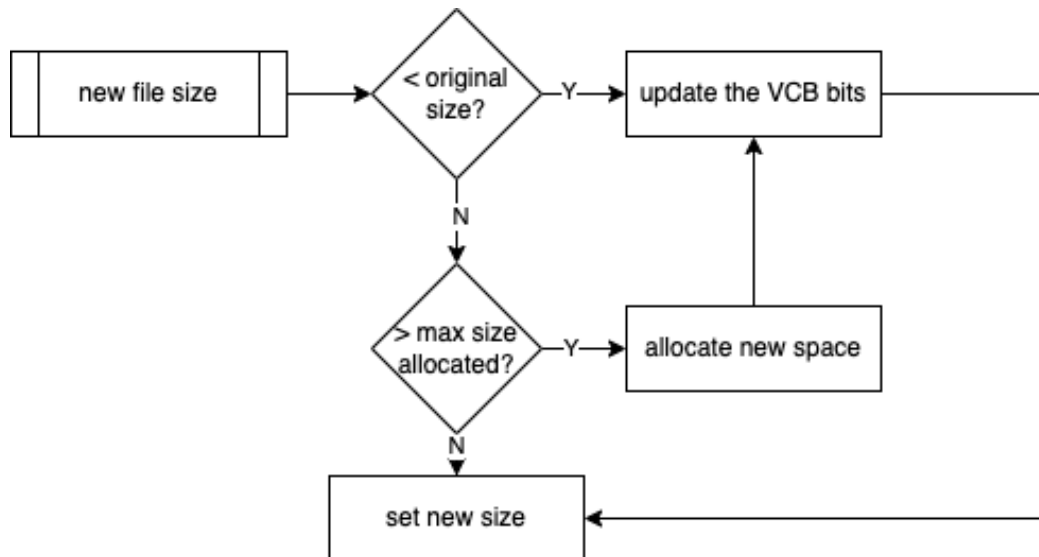
This file write some characters to a file. Notice that the filesystem does not support **append** operation, which means writing is always done by overwriting previous content for a file.

source

When `fs_write()` receives a file pointer, it first checks whether it is valid. If it's not valid, it will print an error and return. Else, there are three scenarios:

1. The new size is smaller than the previous size. We don't need to allocate a new space. However, the valid bits of the file need to be changed (decreased)
2. The new size cannot fit into the previously allocated blocks. This means we have to find a new place to store this file. This process is handled by the `fs_allocate()` function and will be discussed later.
3. The new size is larger than or equal to the previous size, but the previously allocated blocks has enough space to hold the new content. This scenario is quite common in the bonus part, since creating a new file within the directory means the parent directory size need to be modified. In this scenario, there is no reallocation, and the file's VCB bits are remained.

In all of the scenarios, the *size* attribute and the *modify time* attribute of file need to be changed, and the file system time need to be increased.

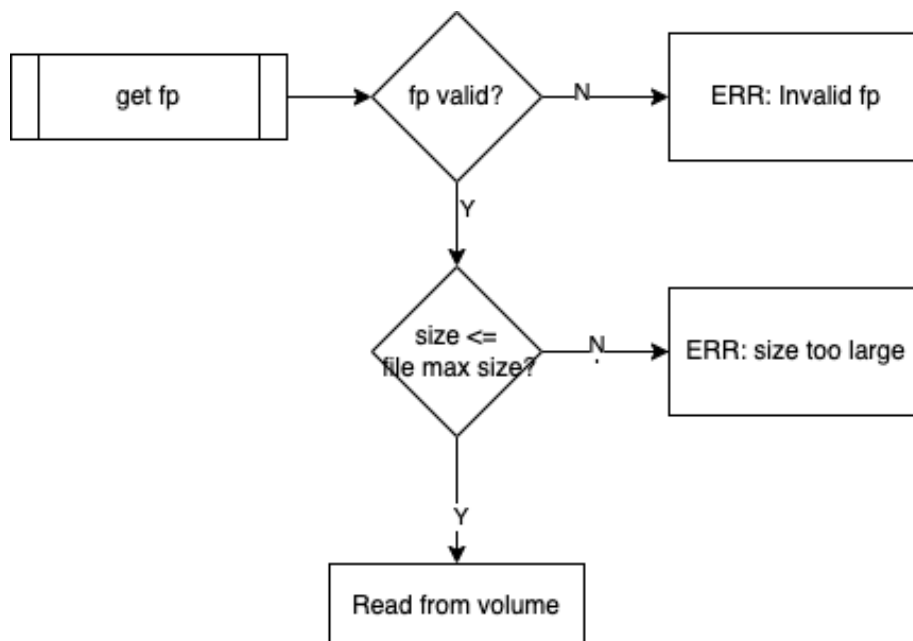


bonus

Aside from updating the *modify time* of the file itself, in the bonus we also have to update the direct parent's *modify time* attribute. Moreover, we also have to update the parent directory's *size* and the file content, because the size of a directory is the number of bits of its contents.

fs_read()

The `fs_read()` is pretty straightforward. It only checks if the `fp` is valid and whether the read size is smaller than the file size. If so, it reads some bytes from the file's base address.



`fs_gsys(LS_D, LS_S)`

source

These two operations list the files in the root directory based on certain criteria. It is achieved through some sorting algorithm. Since calling host functions (like the ones provided in `thrust` library) is not possible in a `_device_` function, I write my own selection sort algorithm. In both **LS_D** and **LS_S**, two temporary arrays are created, one for storing the FCB index and another for storing the corresponding attribute value (modify time/size). These informations are achieved by first traversing all FCBs. Then, I use selection sort to sort both arrays based on the attribute value.

in **LS_D**, since no two files have same modification time, we can directly compare directly. However, if we want to sort based on the file size, we have to compare the creation time of two files to determine the precedence. I do not create the third array for space usage concern. Whenever two files share same size, their creation time is directly read from the volume and compared.

bonus

In the bonus part, instead of listing all the files and directories, we only list the ones present in current working directory. Moreover, a file is differentiated from a folder by checking the *is_dir* bit in the FCB.

`fs_gsys(RM)`

source

In the basic version of the project, removing a file comprises the following steps:

1. Find the FCB index with the requested name
2. Clear the VCB bits
3. Invalidate the FCB
4. Decrement the file number count

Notice that deleting a file acutally preserve the original information in the volume. However, they are not achieveable by the user anymore.

bonus

In the bonus part, aside from the above-mentioned steps, we also have to:

- Prevent the user from using `RM` on a directory
- Pop the file name from the parent directory's content and update the *size* of the parent
- Update the *modify time* of the parent
- Increment the system time

`fs_gsys(RM_RF)` (bonus)

The `RM_RF` can only be used to a directory. It recursively deletes all the files and subdirectories within the directory before ultimately deleting the folder itself. Therefore, it can be viewed as calling `fs_gsys(RM_RF)` and `fs_gsys(RM)` depending on the type of content.

Lastly, just like `fs_gsts(RM)`, after deleting the directory itself, we also have to update the parent directory's content and size and update the modify time, the number of files and the system time.

`fs_gsys(CD)` (bonus)

This function is simply done by updating the global variable of current working directory and incrementing the level depth count.

`fs_gsys(CD_P)` (bonus)

This function is done by retrieving the fp information from the global variable of current working directory, replacing it with its parent's fp, and decrementing the global level depth count.

`fs_gsys(MKDIR)` (bonus)

The `MKDIR` function is actually calling the `fs_open()` to create a new directory, because in the file system's point of view, a directory is simply a file whose content is the files and subdirectories within this directory. Therefore, in `fs_gsys(MKDIR)`:

1. We have to validate this directory's FCB, set the *is_dir* bit and the *parent_fp* bits
2. Set the *name*, *create_time*, *modify_time* attribute of this directory's FCB
3. Update the parent FCB's content

`fs_gsys(PWD)` (bonus)

This function is done by looping indefinitely, starting from the current working directory, we can find its name and parent fp, and its parent's parent..... until the root directory. We can construct an array with size global level depth count - 1, because we don't print the root directory unless we are in it.

e.g: if the current working directory is `/app/soft`, we do not print `//app/soft`, but if we are in root directory, we print `/`. Notice we are traversing to the root directory bottom-up, so we have to print the last element in the array first.

Testcase Output

Case 1:

```

[120090472@node21 source]$ ./main
===sort by modified time===
t.txt
b.txt
===sort by file size===
t.txt          32
b.txt          32
===sort by file size===
t.txt          32
b.txt          12
===sort by modified time===
b.txt
t.txt
===sort by file size===
b.txt          12
[120090472@node21 source]$

```

	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F	Decoded Text
000000	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000010	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0000A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0000B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0000C0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0000D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0000E0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0000F0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0001A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0001B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0001C0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0001D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0001E0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0001F0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000200	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000210	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
000220	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

Case 2:

```
[120090472@node21 source]$ ./main
===sort by modified time===
t.txt
b.txt
===sort by file size===
t.txt      32
b.txt      32
===sort by file size===
t.txt      32
b.txt      12
===sort by modified time===
b.txt
t.txt
===sort by file size===
b.txt      12
===sort by file size===
*ABCDEFGHIJKLMNPQR      33
)ABCDEFGHIJKLMNPQR      32
(ABCDEFGHIJKLMNPQR      31
'ABCDEFGHIJKLMNPQR      30
&ABCDEFGHIJKLMNPQR      29
%ABCDEFGHIJKLMNPQR      28
$ABCDEFGHIJKLMNPQR      27
#ABCDEFGHIJKLMNPQR      26
"ABCDEFGHIJKLMNPQR      25
!ABCDEFGHIJKLMNPQR      24
b.txt      12
===sort by modified time===
*ABCDEFGHIJKLMNPQR
)ABCDEFGHIJKLMNPQR
(ABCDEFGHIJKLMNPQR
'ABCDEFGHIJKLMNPQR
&ABCDEFGHIJKLMNPQR
b.txt
```

	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F	Decoded Text
000000	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000010	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0000A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0000B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0000C0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0000D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0000E0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0000F0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0001A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0001B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0001C0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0001D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0001E0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
0001F0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000200	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000210	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x
000220	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	x x x x x x x x x x x x x x

Case 3:

<omitted same part with previous testcases>

```
===sort by file size===
```

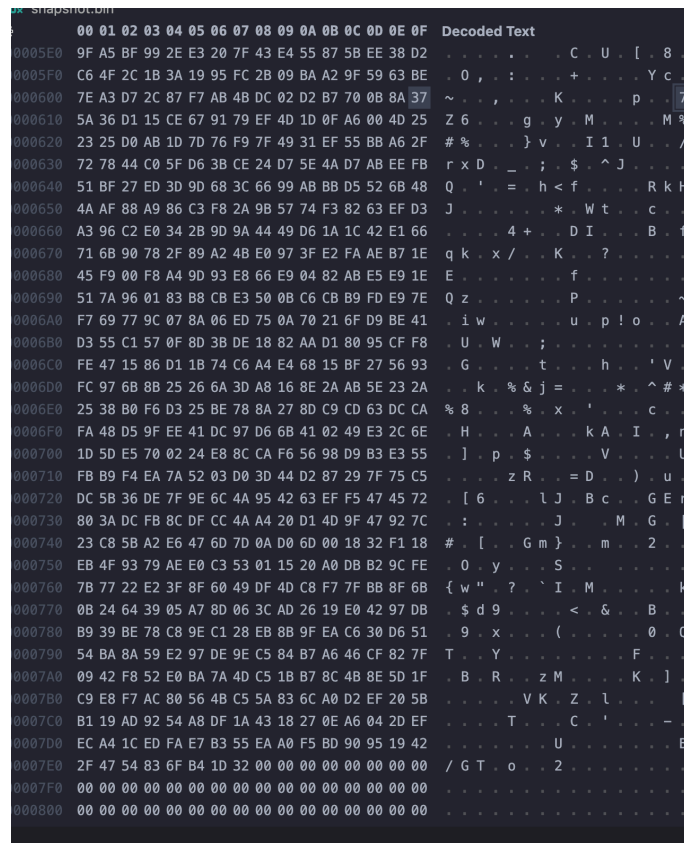
```
EA 1024
~ABCDEFGHJKLM 1024
aa 1024
bb 1024
cc 1024
dd 1024
ee 1024
ff 1024
gg 1024
hh 1024
ii 1024
jj 1024
kk 1024
ll 1024
mm 1024
nn 1024
oo 1024
pp 1024
qq 1024
}ABCDEFGHJKLM 1023
|ABCDEFGHJKLM 1022
{ABCDEFGHJKLM 1021
zABCDEFGHJKLM 1020
yABCDEFGHJKLM 1019
xABCDEFGHJKLM 1018
wABCDEFGHJKLM 1017
vABCDEFGHJKLM 1016
uABCDEFGHJKLM 1015
tABCDEFGHJKLM 1014
sABCDEFGHJKLM 1013
rABCDEFGHJKLM 1012
qABCDEFGHJKLM 1011
pABCDEFGHJKLM 1010
```

< Omitted size from 1009 to 52>

```
MA 51
LA 50
KA 49
JA 48
IA 47
HA 46
GA 45
FA 44
DA 42
CA 41
BA 40
AA 39
@A 38
?A 37
>A 36
=A 35
<A 34
*ABCDEFGHJKLMNQPQR 33
;A 33
)ABCDEFGHJKLMNQPQR 32
:A 32
(ABCDEFGHJKLMNQPQR 31
9A 31
'ABCDEFGHJKLMNQPQR 30
8A 30
&ABCDEFGHJKLMNQPQR 29
7A 29
6A 28
5A 27
4A 26
3A 25
2A 24
b.txt 12
[120090472@node21 source]$
```

snapshot.bin																	Decoded Text
	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F	
000000	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000010	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	6F	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0000A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0000B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0000C0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0000D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0000E0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0000F0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0001A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0001B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0001C0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0001D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0001E0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0001F0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000200	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000210	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
000220	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

0x snapshot.bin																	Decoded Text
	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F	
00003D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
00003E0	00	00	00	00	00	00	00	00	FA	E5	17	F8	A2	DC	72	AF	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
00003F0	4B	A0	28	C0	C3	18	25	EE	E6	9F	F4	44	5A	7E	8E	E9	K . (. . . % . . . D Z ~ . . .
0000400	95	C5	E4	24	E3	74	29	DE	D9	BF	57	FC	9D	CA	AC	E8	. . . \$. t) . . . W . . .
0000410	6B	54	2A	EA	EB	CE	1D	D3	EE	12	97	49	90	A5	B2	A6	k T * I . . .
0000420	68	97	CA	50	8C	73	AE	66	34	07	63	D1	D1	10	3A	BC	k . . P . s . f 4 . c . . . : . .
0000430	64	E3	6B	51	B3	88	A4	22	1A	BB	6A	AA	61	9D	51	CC	d . k Q . . . " . . k . a . Q . . .
0000440	B5	9C	9C	42	10	4C	A8	44	53	0D	95	A4	9C	50	E0	81	. . . B . L . D S . . . P . . .
0000450	B3	CB	D2	67	54	F6	89	ED	B2	74	98	14	92	6A	60	48	. . . g T . . . t . . . j ` h . . .
0000460	07	FD	8A	96	4A	33	DB	1D	BF	F0	41	5D	C0	22	DE	75	. . . J 3 . . . A] . " . u . . .
0000470	ED	31	5C	C1	28	66	AF	5A	DA	C7	ED	EC	B1	4E	35	38	. 1 \ . (f . Z . . . N 5 8 . . .
0000480	CB	3F	CF	95	F2	2B	32	B2	1C	73	10	DD	95	6E	53	03	. ? . . . + 2 . . . s . . . n S . . .
0000490	1F	AF	44	C6	16	F3	21	71	3C	8E	DD	ED	5D	14	27	29	. D . . . ! q < . . .] . ') . . .
00004A0	53	F6	3F	C5	22	71	79	BD	E5	09	1B	FA	F7	ED	7E	17	S . ? . " q y ~ . . .
00004B0	1E	C2	DE	B3	B7	00	A4	F3	8F	83	61	EC	17	88	95	E9 a
00004C0	FE	D4	B0	21	47	2A	5F	AC	33	7A	A7	AA	E8	26	C2	07	. . . ! G * _ . 3 z . . . & . . .
00004D0	E9	A1	BA	21	21	DF	94	30	63	F5	1D	7A	FE	33	64	FD	. . . ! ! . 0 c . . z . 3 d . . .
00004E0	87	15	9F	CE	BE	FE	FA	72	F8	A3	1D	61	49	DF	68	33 r . . . a I . h s . . .
00004F0	81	A3	D3	23	83	E7	53	E6	5E	F0	E0	5D	24	C4	5B	AB	. . . # . . S . ^ . .] \$. [. . .
0000500	DA	7A	FA	19	F8	F5	8B	72	19	A9	D3	63	09	3D	16	0B	. z r . . . c . = . . .
0000510	60	EA	2E	E3	52	81	4A	B0	72	2B	0E	16	EF	E9	42	4A	` . . . R . J . r + . . . B C . . .
0000520	64	BC	64	DD	32	6F	50	4C	98	24	AF	A2	61	45	2D	41	d . d . 2 o P L \$. . a E - A . . .
0000530	AF	5B	25	03	5C	EE	B3	4F	99	42	65	0A	2C	27	54	10	. [% \ . . 0 . B e . , ' T . . .
0000540	E3	38	ED	17	28	3E	63	C0	E2	92	63	44	D7	90	06	88	. 8 . . (> c . . . c D . . .
0000550	6B	2B	0B	C8	9A	BE	18	34	80	FC	3E	2C	25	13	3D	09	k + 4 . . . > , % = . . .
0000560	CA	AA	20	F2	69	03	B4	4C	95	97	10	ED	A8	16	F5	14 i . . L
0000570	42	01	5C	DC	3F	F3	90	40	F1	CF	6C	17	E2	A9	0F	AD	B . \ . ? . . @ . . l . . .
0000580	55	C0	A1	BE	43	D5	8A	D9	6D	9A	47	95	B1	3D	2A	F3	U . . C . . m . G . = * . . .
0000590	BD	86	50	7C	7B	E0	BC	6D	30	2A	04	92	53	A3	C0	28	. . P { . . m 0 * . . S . . .
00005A0	E3	E1	66	28	B7	F0	81	A4	8C	C8	3B	3E	85	65	B1	C2	. . f (. ; > . e . . .
00005B0	6B	81	BE	E6	E1	7C	D3	13	26	57	25	79	7B	E5	A2	5F	k & W % y { . . .
00005C0	C7	88	87	7F	7A	88	A4	07	D0	5F	C4	D5	C4	76	98	AF z . . . _ . . v . . .
00005D0	F7	58	97	59	54	EA	6C	7A	C2	12	73	3E	F7	16	9D	C0	. X . Y T . l z . . s > . . .
00005E0	9F	A5	BF	99	2E	E3	20	7F	43	E4	55	87	5B	EE	38	D2 C . U . [. 8 . . .
00005F0	C6	4F	2C	1B	3A	19	95	FC	2B	09	BA	A2	9F	59	63	BE	. 0 , . : . . + Y c . . .



Case 4:

```
[120090472@node21 source]$ time ./main
triggering gc
===sort by modified time===
1024-block-1023
1024-block-1022
1024-block-1021
1024-block-1020
1024-block-1019
1024-block-1018
1024-block-1017
1024-block-1016
1024-block-1015
1024-block-1014
1024-block-1013
1024-block-1012
1024-block-1011
1024-block-1010
1024-block-1009
1024-block-1008
1024-block-1007
1024-block-1006
1024-block-1005
1024-block-1004
1024-block-1003
1024-block-1002
1024-block-1001
1024-block-1000
1024-block-0999
1024-block-0998
1024-block-0997
1024-block-0996
1024-block-0995
1024-block-0994
1024-block-0993
```

<omitted 1024-block-0992 to 1024-block-0030>


```
1024-block-0029
1024-block-0028
1024-block-0027
1024-block-0026
1024-block-0025
1024-block-0024
1024-block-0023
1024-block-0022
1024-block-0021
1024-block-0020
1024-block-0019
1024-block-0018
1024-block-0017
1024-block-0016
1024-block-0015
1024-block-0014
1024-block-0013
1024-block-0012
1024-block-0011
1024-block-0010
1024-block-0009
1024-block-0008
1024-block-0007
1024-block-0006
1024-block-0005
1024-block-0004
1024-block-0003
1024-block-0002
1024-block-0001
1024-block-0000
```

```
real    0m16.760s
user    0m9.774s
sys     0m6.870s
```

```
▶[120090472@node21 source]$ cmp data.bin snapshot.bin
▶[120090472@node21 source]$ █
```

Bonus

```

[120090472@node21 bonus]$ ./main
===Sort 2 files by modified time===
t.txt      4
b.txt      3
===Sort 2 files by size===
t.txt      32
b.txt      32
===Sort 3 files by modified time===
app        5      d
t.txt      4
b.txt      3
===Sort 3 files by size===
t.txt      32
b.txt      32
app        0      d
===Sort 0 files by size===
===Sort 3 files by size===
a.txt      64
b.txt      32
soft       0      d
===Sort 3 files by modified time===
soft       10     d
b.txt      9
a.txt      7
/app/soft
===Sort 4 files by size===
B.txt      1024
C.txt      1024
D.txt      1024
A.txt      64
===Sort 3 files by size===
a.txt      64
b.txt      32
soft       24     d

```

```

/app
===Sort 3 files by size===
t.txt      32
b.txt      32
app        17     d
===Sort 2 files by size===
a.txt      64
b.txt      32
===Sort 3 files by size===
t.txt      32
b.txt      32
app        12     d
[120090472@node21 bonus]$

```

```
[120090472@node21 source]$ cmp data.bin snapshot.bin  
[120090472@node21 source]$
```

Problems encountered

Revealing the status of the file system

In the debugging process, it is often necessary to figure out the current status of the file system. However, CUDA-GDB is often not very helpful, as I find some variables cannot be printed out. Therefore, I write a function `fs_diagnose()` that prints all metadata of the valid FCBs.

Allocating new space using `fs_allocate()`

This function simulates the allocation of a file that is too large to fit in the previously allocated block, which could happen if:

1. The file is being written for the first time, and previously has not been allocated with any space
2. The file's new size is greater than the previously allocated space

To allocate a new space, we adopt **Next Fit Algorithm**. In this algorithm, we use the global variable `glastblock`, which records the first block after the previous allocation in the VCB. Starting from that position, the program will traverse sequentially, looking for n consecutive free bits in the VCB, where n is the new block size of the file. When no such block is found, a file compression is needed.

Volume compression using `fs_compress()`

We need compression because our file system generates **external fragmentation** during its normal usage. An external fragmentation happens when we delete a file or overwrite a file with less content. This means that although there are enough blocks available to store files, they are not contiguous. One way to eliminate such fragmentation is done through compressing the volume. In this operation:

1. We first sort all the files in ascending order based on their *start block index* attribute, which is readily acquired from the FCBs
2. Starting from the first file, we loop through the sorted file list, moving the next file's start block next to the previous file's end block. This operation involves the following operations:
 1. Update the *start block index* of files
 2. Update the VCB bits
 3. Moving the actual file content in the storage area, which is achieved via the `move_file` function.
3. Finally, we also have to update the glastblock value.

It should be noticed that `fs_compress` does not update the *modify time* of the moved files.

Sorting: Time VS Space

In the original desin, I did not allocate any new space when sorting. I traverse the FCB 1024*1024 times. Each time, I find the maximum value that is smaller than the previous maximum value and print it out. However, such operation significantly reduces the program's speed. Therefore, in later designs, whenever I need to do sorting, I always use the heap space to store arrays and free them when the sorting is done.

Learning Outcome

In this project, I learned how to design a file system. Moreover, I learned to apply the programming philosophy of abstraction. In a file system, from the user point of view, he doesn't need to know how a file is opened/read/written. Moreover, some functions in my design is basically wrapping other functions. For example, the `RM_RF` is essentially recursively calling `RM` to remove files. This project also improves my C++ programming and debugging skills. Although in the end I cannot use the standard library or the thrust library, I still appreciate that the USTFs in the WeChat Group mentioned these libraris. They could be very useful later when I write host functions.